

EECE 7360 Project 4
Subset Sum

Garrett Goode and Daniel Hullihen

April 24, 2017

Chapter 1

Introduction

The subset sum problem (also referred to as the “exact knapsack problem”) is defined below.

Let $A = \{a_1, \dots, a_n\}$ represent some set of integers. Given a sum s , find a subset $A' \subset A$ such that

$$s = \sum_{i=1}^m a'_i, \text{ for } 1 \leq i \leq m.$$

Where m is the size of A' .

In other words, if we are given a list of numbers and some target sum, we want to find the numbers in the list that add up to the target sum. Put as a decision problem, the question would be “Is there a subset A' of A where the sum of the elements of A' is s ?”

In this project, several approaches for solving instances of the subset sum problem were investigated and compared against each other. Benchmarks were developed based on prior work, and background research was done to highlight subproblems of subset sum. Exhaustive and greedy solvers were implemented and evaluated. Linear programming (LP) and integer-linear programming (ILP) models were developed and evaluated as well. Lastly, local search algorithms were implemented, evaluated and compared to all other previous implementations.

Chapter 2

Benchmarking and Instance Generation

Before any benchmark or instances can be generated, it is necessary to understand the prior work done to evaluate implementations of algorithms that solve the subset sum problem. The most popular metric used for determining the difficulty of an instance of the subset sum problem is the “density” of the set in question. The following equation describes the density of a set.

$$d = n / \log_2 \max(a_i)$$

where n is the size of the set in question and $\max(a_i)$ is the largest element in the set.

Previous work has focused on being able to solve a group of instances that had a some density [4]. For example, Radziszowski and Kreher explored efficiently solving low density subset sum problems ($d < 0.7$) [7]. Schnorr and Euchner pointed out “the hardest subset sum problems turn out to be those that have a density that is slightly larger than 1, i.e. a density about $1 + \log_2(n/2)/n$ ” [8]. On top of that, LaMacchia points out that subset sum problems with a density $d < 0.6463\dots$ could be solved in polynomial time by a “lattice oracle” [5], showing that there are existing algorithms designed to work efficiently in certain areas of the instance space (in this case, with a sufficiently low density).

What this boils down to is: much research has focused on subset sum problems with densities of and around 1.0. More recently, when evaluating a GPU implementation of the subset-sum problem, Wan et al. crossed vector sizes of 36 to 54 containing random values in the range $[1, 10^8]$, covering a large range of possible densities [2]. Thus, in order to provide a fair comparison with prior work, we need to create instances with densities on and around 1.0. For breadth, we should focus on sufficiently small and large densities as well, and cover different combinations of set length and maximum value that may yield the same density in order to see if those values may have a disproportionate impact on performance.

2.0.1 Instance Generation

Given the prior work, when it comes to generating benchmarks for evaluating different implementations of the subset sum problem, it is necessary to focus on the density of the instances generated, and the range of densities covered by the overall suite. 152 instances were generated, with an emphasis on instances with densities around 1.0 in order to have a finer granularity of data. To generate each instance, we determined the number of elements needed in the set, as well as the largest value that will exist in the set (determined by the number of bits it represents). For each generated number in the set, b bits of random value 0 or 1 were generated. These bits were then concatenated together to form an integer. Once the list was complete, we randomly chose exactly half of the elements in the list and added them up. This sum is the solution for the instance.

The following table shows the different groups of instances that were generated and some properties that were derived from them. In each row, a range of n and b values were generated. These lists of values were crossed to generate a group.

n (Start:End:Stride)	b (Start:End:Stride)	Count	Min d	Max d	Avg. d
2:10:2	20:26:2	20	0.076	0.5	0.263
16:30:2	15:30:2	64	0.55	2.0	1.09
2:10:2	2:10:2	25	0.2	5.0	1.37
50:100:10	15:30:5	24	1.66	6.66	3.5
90:100:2	2:10:3	18	11.25	50	26.125

Note it is possible to achieve the same density with different values of n and b . In general, we aimed to cover a cross of large and small n with large and small b (e.g. large set with small numbers, small set with large numbers, etc.). Together, the above groups in the table span densities from as low as 0.076 to as high as 50, covering a wider range than what was described in the literature. Given the emphasis on cases around a density of 1.0, we argue this is representative of the instance space and in line with prior work.

Chapter 3

Exhaustive Algorithm

For the first attempt at solving the Subset Sum problem instances we had generated, we chose an exhaustive "brute force" approach. Rather than attempt to target specific subsets of our input set, we cycled iteratively through every possible solution.

Algorithm 1 Exhaustive Algorithm to Solve Subset Sum

```
for subset s in set S do
    sum  $\leftarrow$  0
    for element i in s do
        sum + = element
    end for
    if sum == targetSum || timeOutReached then
        break
    end if
end for
```

Since an element in the input set can be "in" or "out" of the solution set, there are 2^n possible subsets. Therefore, the possible subsets can all be represented by an n-digit binary number. In our implementation of a data representation of a subset sum instance, we represented a solution as an array with the same number of elements as the input set, with each entry set to "INCLUDED" or "EXCLUDED". This representation allowed us to emulate adding 1 to a binary number, starting with every entry set to EXCLUDED and ending when every entry was set to INCLUDED.

Given the above, it is easy to characterize the run time of our exhaustive algorithm. There is an initiation step, and the 2^n possible solutions are looped through in the worst case. Additionally, generating the current sum of the solution set also requires cycling through all n of its elements. This yields the expression below.

$$O(n + n * (2^n)) = O(2^n)$$

As expected, the worst cases of Subset Sum cannot be solved exhaustively in polynomial time.

3.1 Results

The results are presented in the figure below. The vertical axis represents the input size and the horizontal axis the word length of the elements in bits.

Notice that the majority of the instances were solved very quickly. And there were more still that could be solved in less than 1 minute, namely up to an instance size of 30, depending on the max bit-width of the elements in the set. The case with a bit-width of 27, for example, took longer than a minute. The largest instance set that could be solved in 1 minute or less was 94, while the largest size that could be solved in less than 10 minutes was 100. Even then, there were instances between these two cases that could not be solved in 10 minutes.

	2	4	5	6	8	10	15	17	19	20	21	22	23	24	25	26	27	29	30
2	0	0		0	0	0				0		0		0		0			
4	0	0		0	0	0				0		0		0		0			
6	0	0		0	0	0				0		0		0		0			
8	0	0		0	0	0				0		0		0		0			
10	0	0		0	0	0				0		0		0		0			
16							0	0	0		0		0		0		0	0	
18							0	0	0		0		0		0		0	0	
20							0	0	0		0		0		0		0	0	
22							0	0	0		0		0		0		0	0	
24							0	0	1		0		2		0		1	0	
26							0	0	0		1		3		2		6	4	
28							0	0	0		3		6		1		20	27	
30							0	0	1		0		13		34		107	55	
50							402				5				424				
60							600			600					600				
70							600			600					600				
80							600			600					600				
90	600		600		600		600			600					600				29
92	600		600		600														375
94	600		600		600														26
96	600		600		600														600
98	600		600		600														600
100	600		600		600		600			600					600				210

Figure 3.1: Run time for various input size and word length combinations

As expected, the exhaustive "brute force" approach largely failed to solve instances where $n > 50$. This is not universally true however, as the algorithm did manage to solve a few large instances of the largest word size we tested. From this we can conclude that word size does not have a large impact on the run time of the algorithm and that brute force can be viable for even some large instances. However speaking more generally an exhaustive approach is really only viable for smaller instances of subset sum, per both the complexity analysis in the preceeding section and the data presented in the last figure.

3.2 Conclusion

The exhaustive algorithm is a simple algorithm to implement, but can quickly approach its timeout. By definition every single instance will be considered unless a timeout mechanism is used to end the solver early. The worst-case time complexity is not realized in our experiments until the instance size was greater than 60 elements. Even at 50 elements the algorithm began to show signs of struggle. In general, the density of the instance did not impact the run-time of this algorithm; instead, the algorithm was largely impacted by the instance size.

Chapter 4

Survey of Complexity Landscape

In this chapter, we explore the complexity landscape around the subset sum problem, including subproblems to subset sum, as well as problems of which subset sum is itself a subproblem.

4.1 Natural Subproblems of Subset Sum

The subset sum problem can be solved in pseudo-polynomial time, which means there are subproblems to subset sum that can be solved in polynomial time, but the “hardest” subproblem in subset sum is nevertheless NP-Complete. Due to this property, it is possible to detect some subproblems to subset sum and apply a known algorithm that solves the problem in polynomial time. Some algorithms used for solving subset sum use a divide-and-conquer approach and identify these subproblems. Subproblems to subset sum, for example, can have all numbers that have some special property, or the set of numbers itself has some property that can be exploited. This section explores some of the natural subproblems to subset sum.

4.1.1 Sets with a low density

One metric for describing an instance of subset sum is the density of the set. The following equation describes the density of a set.

$$d = n / \log_2 \max(a_i)$$

where n is the size of the set in question and $\max(a_i)$ is the largest element in the set. Lagarias and Odlyzko found that, for sets that have a density less than 0.645, their “Algorithm SV” can solve the instance in polynomial time [4]. One of the steps in their algorithm involves taking set and perform a transformation in order to apply an algorithm developed by Lenstra, Lenstra, and Lovasz, which has a time complexity of $O(n^{12} + n^9(\log|f|^3))$ [6]. The transformed problem is one where a short vector e must be found within an integer lattice $L = L(a, M)$

4.1.2 Sets where all the numbers are coprime to m

This approach (as well as the prior approach) focus on sets that are finite cyclic groups, defined as $\mathbb{Z}_m = 0, 1, \dots, m-1$ of order m . Koiliaris and Xu argue that, given a set $S \subseteq U(\mathbb{Z})_m$, which describes a set of integers that are coprime to some integer m , finding the set of all subset sums can be done with a time complexity of $O(\min(\sqrt{nm}, m^{5/4}) \log m \log n)$ [3]. Note that an integer x is coprime to m if the greatest common denominator of the two values is 1. As Koiliaris explains, the method behind the speed-up with this approach is the ability to partition the set into different subsets such that “every such subset is contained in an arithmetic progression of the form $x, 2x, \dots, lx$ ”. With this, one can calculate the subset sums very quickly by scaling l accordingly. These sums can then be added together. This is only doable if m is a prime number, or if all the numbers are relative prime to m .

4.1.3 Set is a subset of \mathbb{Z}_m

Koiliaris and Xu are able to take the previous case and take it a step further, developing an algorithm that finds all of the subsets of a set S in with a time complexity of $O(\min(\sqrt{nm}, m^{5/4}) \log^2 m)$ [3]. This is for sets $S \subseteq \mathbb{Z}_m$, which is slightly different from the previous case in that the elements of the set no longer have to be coprime to some integer m . Their algorithm involves recursively computing a partial subset sum as they work across subsets of the original set.

4.1.4 Set size $m \geq l^{1/\alpha}$ and elements are distinct

Chaimovich, Freiman and Galil found that, if they took the subset sum problem and restricted it such that $\max_i |a_i| \leq l \leq m^\alpha$ (where l is some bound and m is the number of variables in the set), then the instance could be solved with a time complexity of $O(lm^2)$ [1]. One unique feature of their algorithm is that supports large sets (previous algorithms could only handle moderately large sets). The main approach in their algorithm is, assuming A^* is defined as $S_b | B \subseteq A$, and S_b is defined as $\sum_{a_i \in B} a_i$, characterize A^* as “a small collection of arithmetic progressions.”

4.2 Problems of which Subset Sum is a Natural Subproblem

We were unable to identify any problems of which Subset sum is a natural subproblem. However rather than ignore this area of investigation entirely, we chose to explore similar problems that are close neighbors via transformation in the landscape of common NP-complete problems.

4.2.1 Satisfiability

While Satisfiability (SAT) is a few transformations away from Subset Sum on the complexity landscape we have presented, establishing this link between subset sum and SAT is crucial as almost all NP completeness proofs are derived from Cook’s original proof for SAT and some transformations. The problem is defined below.

Given a set of U variables, collection of C clauses over U , is there a satisfying truth assignment for C ?

As proven by Cook in the aforementioned paper, SAT is NP-complete.

4.2.2 3-Satisfiability

Transformed from SAT, 3-Satisfiability (3SAT) is a special case of SAT where each logical clause can have only exactly three variables. The problem is formally defined as follows.

Given a set of U variables, collection of C clauses over U such that $|c| = 3$, is there a satisfying truth assignment for C ?

Despite this restriction on the domain of problem instances when compared to SAT, 3SAT remains NP-complete.

4.2.3 3-Dimensional Matching

Transformed from 3SAT, the 3-dimensional matching problem deals with finding a matching in a three-dimensional graph. The formal definition is given below.

Given a set $M \subset W \times X \times Y$, where W , X , and Y are disjoint sets with the same number of elements q , does M contain a matching $M' \subset M$ such that $|M'| = q$ and no two elements are the same?

Like the previous problems discussed in this section, 3DM remains NP-complete in all but a few ideal scenarios. Interestingly enough, 2-Dimensional Matching (2DM) which is not discussed at length in this section, is tractable.

4.2.4 Partition

The Partition problem is defined below.

Let $A = \{a_1, \dots, a_n\}$ represent some set of positive integers. Does there exist a subset $A' \subset A$ such that

$$\sum_{a \in A'} s(a) = \sum_{a \in A - A'} s(a)?$$

In other words, is there a subset of the given set with sum equal to the sum of the members of the original set not included in the subset?

Once again, the partition problem is known to be NP-Complete. Of all of the problems discussed in this section, the partition problem is likely the closest to the Subset Sum problem. However it is not a natural subproblem of Subset Sum as no target is given as input. Instead, the target is implied as the half of the sum of the input set.

One can see clearly how simple a transformation from an instance of Partition to an instance of Subset Sum would be however, simply by summing the input set and setting the target total to half of the calculated sum.

4.2.5 Knapsack

The Knapsack problem is also a transformation from the Partition problem (like Subset Sum). As one might expect, the problems are therefore very similar in nature. The Knapsack problem is defined as follows.

Given a finite set U , with each $u \in U$ having an associated size $s(u)$ and value $v(u)$, is there a subset $U' \subset U$ such that the sum of the $s(u')$ for all $u' \in U'$ is less than a given integer K while the sum of $v(u')$ for all $u' \in U'$ is greater than a given integer H ?

Like all of the other problems in this section, knapsack is NP-complete. There are variations of the Knapsack problem that are tractable, but they are not discussed in this paper as they are less related to the Subset Sum problem in question. Subset Sum itself is sometimes referred to as the exact knapsack problem.

4.3 Summary

The chart in Figure 1 on the following page maps out all problems discussed in the preceding sections. Note that all subproblems are tractable, while all problems related by transformation are believed to be intractable.

4.4 Conclusion

Overall the investigation into the subproblems and closely related NPC problems was very helpful in getting a clearer vision of the nature of the Subset Sum problem we will continue to investigate. We had some difficulty identifying any problems that subset sum was a natural subproblem of, but in the end that seemed to be appropriate as the subset sum problem itself is very general. It is quite simple to see how when bounds and restrictions are introduced to the original subset sum problem statement we arrive at various other important or interesting subproblems, but it is difficult to relax the input domain of a problem that is already so broad and nonspecific in nature. All in all we are eager to apply the information we discovered in this investigation to aid in our understanding of future solution methods for the Subset Sum problem.

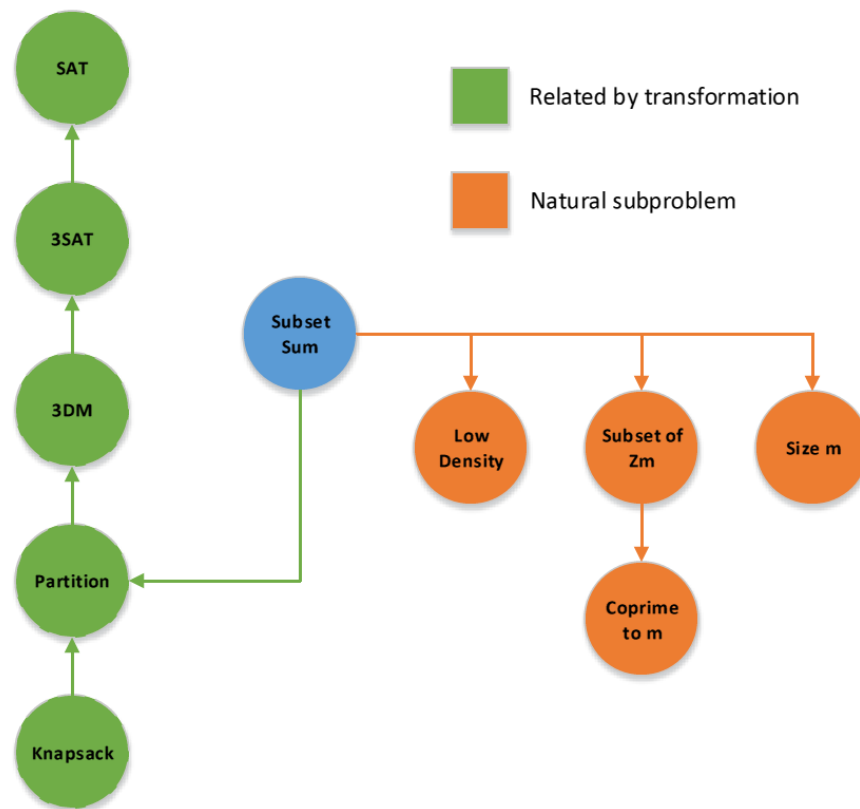


Figure 4.1: Complexity landscape for Subset Sum

Chapter 5

Greedy Algorithm

In this chapter, we review an implementation of a greedy algorithm for the subset sum problem. This implementation was run against a suite of instances, and the performance of this algorithm is described.

5.1 Greedy Algorithm Implementation

The implementation of the greedy algorithm that was used for this project is relatively simple. Below is the pseudocode for the algorithm.

Algorithm 2 Greedy Algorithm to Solve Subset Sum

```
runningSum  $\leftarrow$  0
subset  $\leftarrow$  [] {Initialize subset to an empty list}
for integer i in set S do
    if runningSum + i  $\leq$  targetSum then
        runningSum  $\leftarrow$  runningSum + i
        append i to subset
    end if
end for
if runningSum == targetSum then
    return subset
else
    return NULL
end if
```

We walk through each element of the provided instance, adding numbers so long as the running sum has not exceeded the target sum. At the end, we check to see if the running sum matches the target sum; if it does, we have successfully solved the instance and return the subset of integers. Otherwise we have not solved the instance and return NULL.

One characteristic of a greedy algorithm is that, once we decide to include an element to the sum, we never undo the decision. In the case of the subset sum problem, this can easily cause the algorithm to not find the correct list of elements to use for the sum since not all combinations of integers in a list are necessarily going to add up to the target sum. Because of this, a margin of error is introduced where the algorithm may fall short of the target sum. The upside of this algorithm is the time complexity improvement: the algorithm visits each element of the set only once, giving the algorithm a tightly bound time complexity of $O(n)$. This algorithm may be useful for applications that only need a subset that is within some percent of the target sum, especially given the improved time complexity over the brute-force solution.

5.2 Trivial Case Where the Greedy Algorithm Fails

Due to the nature of the greedy algorithm explained in the previous section, we can identify a trivial case where the algorithm will fail. Take for example the set of integers S below where the target sum is 150.

$$S = \{49, 100, 50\}$$

The greedy algorithm will start from the first element and work its way through the list, adding numbers so long as the running sum does not go beyond the target sum of 150. In this case, the algorithm will accept the integers 49 and 100, resulting in a sum of 149. When it encounters 50, it decides not to add this number since it would bring the running sum over the target sum. The algorithm has then completed processing the set, but it has failed to find a solution despite the fact one actually exists.

5.3 Results

The results are presented in Figure 6.4 below. The vertical axis represents the input size and the horizontal axis the word length of the elements in bits.

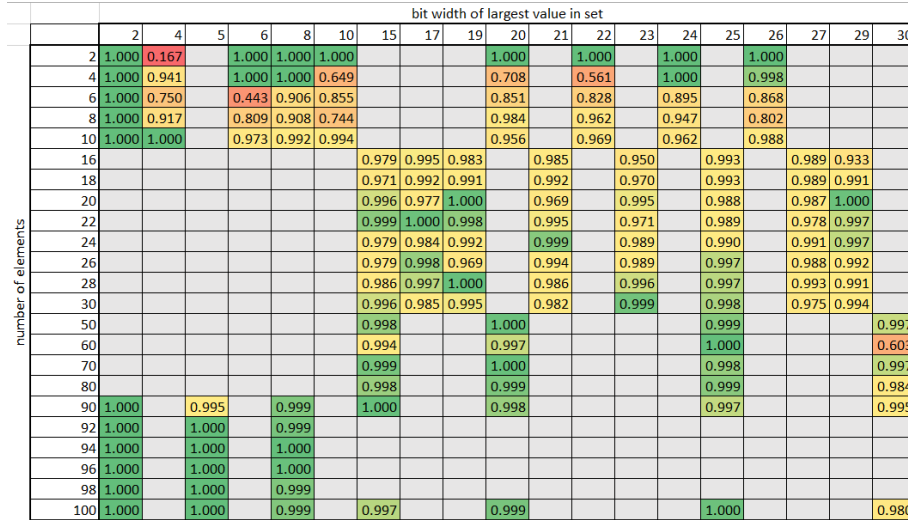


Figure 5.1: Margin of error for various input size and word length combinations for the greedy algorithm

Each cell in the above figure is color-coded based on the margin of error the greedy algorithm had for the given instance. A value of 1 means the algorithm successfully found a subset whose sum of elements matched the target sum. Varying yellow, orange and red elements are cases where the algorithm was more progressively off from the target sum, with red being the most off/greatest error.

Of the 151 instances that were tested, the greedy algorithm was able to correctly solve 28 of them, yielding an 18.5% success rate. Compared to the exhaustive algorithm, which had a 76% success rate, the greedy algorithm is much less successful despite its lower time complexity. As a reminder, Figure 6.5 shows the run-times for the exhaustive algorithm.

Looking at the results described in Figure 6.4, the greedy algorithm was successful when there were sufficiently few numbers in the set, as well large sets with relatively small numbers. For example, the majority of the instances with 90+ elements and a max bit width ≤ 10 passed. This may be due to the fact that, given enough sufficiently small numbers, the algorithm can work at a small enough granularity and have a better chance to come across a total set of numbers that can add up to the target. However, as the size of the max value increases, the margin begins to decrease; the algorithm is less likely to hit the target sum.

The instance where the greedy algorithm had the worst margin was the one with 2 elements and a max bit width of 4. This is most likely due to chance. With fewer numbers, there are fewer chances to find other

	2	4	5	6	8	10	15	17	19	20	21	22	23	24	25	26	27	29	30
2	0	0		0	0	0				0		0		0		0			
4	0	0		0	0	0				0		0		0		0			
6	0	0		0	0	0				0		0		0		0			
8	0	0		0	0	0				0		0		0		0			
10	0	0		0	0	0				0		0		0		0			
16							0	0	0		0		0		0		0	0	
18							0	0	0		0		0		0		0	0	
20							0	0	0		0		0		0		0	0	
22							0	0	0		0		0		0		0	0	
24							0	0	1		0		2		0		1	0	
26							0	0	0		1		3		2		6	4	
28							0	0	0		3		6		1		20	27	
30							0	0	1		0		13		34		107	55	
50							402			5					424				
60							600			600					600				
70							600			600					600				
80							600			600					600				
90	600		600		600		600			600					600				29
92	600		600		600														375
94	600		600		600														26
96	600		600		600														600
98	600		600		600														600
100	600		600		600		600			600					600				210

Figure 5.2: Run time for various input size and word length combinations for the exhaustive algorithm

numbers that can add up to the target sum. It is also a function of the distribution of numbers in the set. If there is a large distribution of integers and the target sum actually needs the larger value, but the larger value is at the end of the set, then the algorithm is more likely to not include it and end up with a large margin of error. One way to possibly correct for this case is by sorting the integers from largest to smallest before processing them.

There appears to be a sufficiently large area where the margin of error is within a few percent of the target sum, but the algorithm nevertheless fails. This appears to happen with medium-sized sets (i.e. sets with more than just a couple elements, but do not have a large number either). This is namely from set sizes of about 6 to 50. The greedy algorithm is unable to correctly solve the majority of these instances. This may be due to the fact that there are enough integers where the algorithm can accidentally add one that will actually never allow the running sum to equal the target sum. Had there been more integers, there would be more opportunities for the algorithm to find some other integer that could still allow the algorithm to achieve the target sum.

In general, the amount of margin exhibited by the greedy algorithm seems to come down to how many subsets actually do exist within the set that satisfy the target sum, which depends on factors such as the distribution of integers values in the set and whether a value repeats itself. For example, for sets with several large and small numbers that are similar to each other, there may be more subsets that satisfy a given target sum compared to a set with a more even distribution of integers. Granted, this also depends on the target sum itself. But if there are many subsets that satisfy the target sum, then the greedy algorithm has a better chance of solving a given instance.

Chapter 6

ILP and LP Modeling

In this chapter, we review LP and ILP techniques. An ILP model was developed. The results of running our ILP model against a suite of instances is described.

6.1 ILP Formulation

The implementation of the ILP formulation we utilized is given by the following AMPL model pseudo-code.

Algorithm 3 AMPL Model for Subset Sum

```
values  $\leftarrow$  {input set}  
X  $\leftarrow$  [] {empty binary array}
```

Ensure:

values[*i*] * *X*[*i*] is maximal

Require:

sum(*values*[*i*] * *X*[*i*]) = target

First parameters are created to store the input data for the problem instance, the set of integers and the target. A second array of binary values is created to track which members of the input set will form the subset that represents the solution.

We elected to maximize the sum of the subset as our objective function, so that even in a situation with no solution we would still be able to achieve the best possible value. This was also the way our greedy algorithm behaved, so it would make it easier to compare the results.

Lastly, our sole condition guaranteed that the sum of the subset would have to be equal to the target. This way, optimal solutions would always stop at the target, despite trying to "maximize" the sum.

6.2 LP Lower Bound

Unsurprisingly, the LP lower bounds we calculated essentially match the ILP results discussed later in the Results portion of the report. This holds with our expectation, as since subset sum seeks an exact value. One notable difference was the LP models ran much faster than their ILP counterparts, as evidenced by the figures below.

6.3 Results

The results are presented in Figure 6.3 below. The vertical axis represents the input size and the horizontal axis the word length of the elements in bits.

		Bit width																																
		2	4	5	6	8	10	15	17	19	20	21	22	23	24	25	26	27	29	30														
Input size	2	0.003733	0.004053				0.003733	0.003309				0.003551	0.003326		0.00353		0.003733																	
	4	0.003564	0.003769			0.00361		0.00348	0.003825			0.003373		0.003391		0.003739																		
	6	0.003623	0.003849					0.003805	0.003862	0.003805			0.003529	0.003689		0.003812		0.003554																
	8	0.003707	0.003534			0.003704	0.003514	0.003439				0.003676	0.003446		0.003697		0.003459																	
	10	0.00355	0.003653				0.003658	0.003729	0.003533			0.003616		0.003771		0.003875		0.00382																
	18							0.004488	0.003615	0.003728			0.003557		0.004179		0.004216		0.00366											0.004369	0.003688			
	18							0.003721	0.003731	0.003865			0.004061				0.00388		0.00389											0.00367				
	22							0.003509	0.003844	0.004084			0.004032		0.004056		0.003936		0.00369	0.003629														
	22							0.003784	0.003628	0.003768			0.003876		0.003494		0.003605		0.003575	0.00353														
	24							0.003683	0.004002	0.003727			0.003491		0.003839		0.003619		0.003323	0.003892														
	28							0.003624	0.004089	0.003784			0.003546		0.003626		0.004048		0.003666	0.003517														
	28							0.004137	0.003854	0.003676			0.003664		0.004005		0.003784		0.003897	0.00406														
	30							0.00385	0.003837	0.003755			0.003627		0.003675		0.003707		0.00366	0.003505														
	50							0.003758									0.003782																	0.003795
	60							0.00382					0.003702				0.003567																	0.003888
	70							0.003879					0.004086				0.003538																	0.003664
	80							0.003887					0.003989				0.003814																	0.003781
	90	0.003442						0.003795		0.003616			0.003724				0.003901																	0.003966
	92	0.003598						0.003867		0.004026							0.003957																	0.004054
	94	0.003645						0.003811		0.003908							0.003901																	0.004006
96	0.00393						0.003544		0.003666							0.003666																	0.003781	
98	0.003614						0.004063		0.003872							0.003872																	0.004006	
100	0.003707						0.003952		0.00365		0.003749		0.003737			0.003564																	0.003859	

Figure 6.1: Runtimes for LP subset sum models

[illegible]

Figure 6.2: Runtimes for ILP subset sum models

Each cell in the above figure is color-coded based on the margin of error the ILP model had in its final solution. This figure is presented mostly to contrast with the other figures in the section, as it is clear the ILP model was able to produce an accurate solution in every instance. Comparing this result with the same figure for the greedy algorithm below, it is abundantly clear that the ILP approach is vastly superior in terms of accuracy. Note that the table in Figure 6.4 uses an accuracy rating that is the inverse of the ILP table - a rating of 1 is a perfect solution and 0 is instead the worst.

Of the 151 instances that were tested, the greedy algorithm was able to correctly solve 28 of them, yielding an 18.5% success rate. Compared to the ILP formulation, which had a 100% success rate, the greedy algorithm is much less successful despite its lower time complexity. As a reminder, Figure 6.5 shows the run-times for the exhaustive algorithm.

In general, the amount of margin exhibited by the greedy algorithm seems to come down to how many subsets actually do exist within the set that satisfy the target sum, which depends on factors such as the distribution of integers values in the set and whether a value repeats itself. For example, for sets with several large and small numbers that are similar to each other, there may be more subsets that satisfy a given target sum compared to a set with a more even distribution of integers. Granted, this also depends on the target sum itself. But if there are many subsets that satisfy the target sum, then the greedy algorithm has a better chance of solving a given instance.

When compared to the exhaustive algorithm, the accuracy advantage of the ILP approach is similar to that of the greedy algorithm. However, when the timing results of the ILP algorithm in Figure 6.2 are compared to that of the exhaustive algorithm, it is actually the much more primitive exhaustive algorithm that has the advantage in solving smaller or less complex instances. This result was surprising, as all other findings seemed to indicate that ILP was more or less a 'magic bullet' for solving our subset sum instances.

		Bit width																												
		2	4	5	6	8	10	15	17	19	20	21	22	23	24	25	26	27	29	30										
Input size	2	0	0		0	0	0				0		0		0		0													
	4	0	0		0	0	0				0		0		0		0													
	6	0	0		0	0	0				0		0		0		0													
	8	0	0		0	0	0				0		0		0		0													
	10	0	0		0	0	0				0		0		0		0													
	16							0	0	0		0		0		0		0		0		0		0		0		0		
	18							0	0	0		0		0		0		0		0		0		0		0		0		
	20							0	0	0		0		0		0		0		0		0		0		0		0		
	22							0	0	0		0		0		0		0		0		0		0		0		0		
	24							0	0	0		0		0		0		0		0		0		0		0		0		
	26							0	0	0		0		0		0		0		0		0		0		0		0		
	28							0	0	0		0		0		0		0		0		0		0		0		0		
	30							0	0	0		0		0		0		0		0		0		0		0		0		
	50										0																			0
	60																													0
	70																													0
	80																													0
	90	0		0		0		0			0						0													0
	92	0		0		0		0																						0
	94	0		0		0		0																						0
	96	0		0		0		0																						0
	98	0		0		0		0																						0
	100	0		0		0		0			0						0													0

Figure 6.3: Margin of error for various input size and word length combinations for ILP

		bit width of largest value in set																												
		2	4	5	6	8	10	15	17	19	20	21	22	23	24	25	26	27	29	30										
number of elements	2	1.000	0.167		1.000	1.000	1.000				1.000		1.000		1.000		1.000													
	4	1.000	0.941		1.000	1.000	0.649				0.708		0.561		1.000		0.998													
	6	1.000	0.750		0.443	0.906	0.855				0.851		0.828		0.895		0.868													
	8	1.000	0.917		0.809	0.908	0.744				0.984		0.962		0.947		0.802													
	10	1.000	1.000		0.973	0.992	0.994				0.956		0.969		0.962		0.988													
	16							0.979	0.995	0.983		0.985		0.950		0.993		0.989	0.933											
	18							0.971	0.992	0.991		0.992		0.970		0.993		0.989	0.991											
	20							0.996	0.977	1.000		0.969		0.995		0.988		0.987	1.000											
	22							0.999	1.000	0.998		0.995		0.971		0.989		0.978	0.997											
	24							0.979	0.984	0.992		0.999		0.989		0.990		0.991	0.997											
	26							0.979	0.998	0.969		0.994		0.989		0.997		0.988	0.992											
	28							0.986	0.997	1.000		0.986		0.996		0.997		0.993	0.991											
	30							0.996	0.985	0.995		0.982		0.999		0.998		0.975	0.994											
	50							0.998			1.000					0.999			0.997									0.997		
	60							0.994			0.997					1.000			0.603									0.997		
	70							0.999			1.000					0.998			0.984									0.984		
80							0.998			0.999					0.999			0.995									0.995			
90	1.000		0.995		0.999		1.000			0.998					0.997															
92	1.000		1.000		0.999																									
94	1.000		1.000		1.000																									
96	1.000		1.000		1.000																									
98	1.000		1.000		0.999																									
100	1.000		1.000		0.999		0.997			0.999					1.000			0.980										0.980		

Figure 6.4: Margin of error for various input size and word length combinations for the greedy algorithm

6.4 Conclusion

Overall it seems that ILP is the best approach we have tested to date for solving subset sum instances, in a general sense. More specifically, ILP excels at solving more complex instances that greedy or exhaustive approaches cannot solve in a reasonable amount of time. However, due to its sophisticated approach to generating a solution, it seems to be excessively costly for solving simpler instances. In these instances, one can reach an optimal solution much faster by using a simpler approach. This illustrates that not only is there no perfect method to solve any optimization problem, but it seems that often there is no perfect method for solving every instance of a single optimization problem.

	2	4	5	6	8	10	15	17	19	20	21	22	23	24	25	26	27	29	30
2	0	0			0	0	0				0	0		0		0			
4	0	0			0	0	0				0	0		0		0			
6	0	0			0	0	0				0	0		0		0			
8	0	0			0	0	0				0	0		0		0			
10	0	0			0	0	0				0	0		0		0			
16								0	0	0		0		0		0		0	0
18								0	0	0		0		0		0		0	0
20								0	0	0		0		0		0		0	0
22								0	0	0		0		0		0		0	0
24								0	0	1		0		2		0		1	0
26								0	0	0		1		3		2		6	4
28								0	0	0		3		6		1		20	27
30								0	0	1		0		13		34		107	55
50								402			5					424			
60								600			600					600			
70								600			600					600			
80								600			600					600			
90	600		600		600		600			600					600				29
92	600		600		600														375
94	600		600		600														26
96	600		600		600														600
98	600		600		600														600
100	600		600		600		600			600					600				210

Figure 6.5: Run time for various input size and word length combinations for the exhaustive algorithm

Chapter 7

Local Search Algorithms

We did local search stuff.

Chapter 8

Conclusion

Nothing yet dawg.

Bibliography

- [1] Mark Chaimovich and Gregory Freiman. Solving dense subset-sum problems by using analytical number theory. *Journal of Complexity*, 1989.
- [2] Lanjun Wan et al. Gpu implementation of a parallel two-list algorithm for the subset-sum problem. *Concurrency and Computation: Practice and Experience*, 27.1:119–145, 2015.
- [3] Konstantinos Koiliaris and Chao Xu. A faster pseudopolynomial time algorithm for subset sum. *Proceedings of the Twenty-Eighth Annual ACM-SIAM Symposium on Discrete Algorithms*, 2016.
- [4] Jeffrey C. Lagarias and Andrew M. Odlyzko. Solving low-density subset sum problems. *Journal of the ACM (JACM)*, 32(1):229–246, 1985.
- [5] Brian A. LaMacchia. Basis reduction algorithms and subset sum problems. Master’s thesis, Massachusetts Institute of Technology, 1991.
- [6] Hendrik Willem Lenstra Lenstra, Arjen Klaas and Laszlo Lovasz. Factoring polynomials with rational coefficients. *Mathematische Annalen*, 261(4):515–534, 1982.
- [7] Stanislaw Radziszowski and Donald Kreher. Solving subset sum problems with the l3 algorithm. *The Charles Babbage Research Centre: The Journal of Combinatorial Mathematics and Combinatorial Computing*, 3, 1988.
- [8] C. P. Schnorr and M. Euchner. Lattice basis reduction: Improved practical algorithms and solving subset sum problems. *Mathematical programming*, 66.1-3:181–199, 1994.